

Data Wrangling with R

2020-01-13

Contents

Introduction	1
R for Data Science	2
Online documentation for the tidyverse	2
Data in R	3
Kinds of variables	3
Vectors, Lists, Data Frames, and Tibbles.	4
The Tidyverse	7
Re-shaping tibbles: gather and spread	14
Graphing Data	18

Introduction

This is a brief introduction to data types, data structures, and some of the functions and packages that we will use to manipulate data in the labs.

This document was generated using RMarkdown and if you want to see the source code for it, you can clone the GitHub repository <https://github.com/gilligan-ees-3310-2020/lab-02-documentation>.

There is a lot more, and two particular resources that I would recommend to you are available free on the web.

R for Data Science

The first is the book, *R for Data Science*, by Hadley Wickham who wrote most of the packages in the `tidyverse` collection. You can buy a print version of the book from all the usual online sources, but Wickham has also posted the full text on the web at <http://r4ds.had.co.nz/> to make it available for free. (Also, he wrote the whole book in RMarkdown, and if you're curious you can get the RMarkdown from <https://github.com/hadley/r4ds>).

The key parts of the book, from the perspective of the labs for this course, are Chapter 4: “Workflow Basics,” which presents a brief overview of R and how to program with it; Chapter 5: “Data Transformation,” which explains the functions I discuss in the first part of this handout about tibbles and the manipulating them with functions like `select`, `filter`, `mutate`, and `summarize`; and Chapter 3: “Data Visualization,” which describes using the `ggplot2` package to make graphs and charts of your data.

If you are interested in learning more about R, Section II of the book discusses the different data types that R uses in detail (tibbles, character data (or strings), factors, dates, and times). Section III discusses programming, and section IV discusses statistical modeling (i.e., fitting functions to data). Section V discusses RMarkdown and all the different ways you can use it to communicate about your analysis with other people.

The book is an excellent introduction to data analysis with R. I have recommended it to many people who did not previously have experience working with programming or R and they found it a very accessible, useful, and user-friendly introduction.

Online documentation for the tidyverse

R for Data Science is a great introduction to the concepts behind the `tidyverse` collection of packages and functions for R, but what should you do when you already understand that big picture and just want to know how to do a specific task? For that, the online documentation for the `tidyverse` is very useful and you can find it at <http://www.tidyverse.org/packages/>.

This page has links to the documentation for all the major `tidyverse` packages: `ggplot2` for making graphics, `dplyr` and `tidyr` for working with `data.frames` and tibbles, `readr` for reading in data from text files on the disk and `readxl` for reading data from Excel spreadsheets, and many more packages that we will not be using in these labs.

The documents give lots of examples showing what the functions do and explanations of how to do many common tasks. Especially for `ggplot2`, it can be very useful to look at the graphs in the examples to find something that looks like what you're trying to do and then seeing the code that made that happen.

Data in R

Kinds of variables

R is capable of analyzing many different kinds of data. Some of the most important kinds of data that we may work with are:

- **Integer data**, which represents discrete quantities, such as counting events or objects.
- **Real number data**, which represents quantities that can have fractional values. Most of the data we will work with in this course, such as temperatures, altitudes, amounts of rain-fall, and so forth, will be real-number data. This kind of data is also referred to frequently as “floating point” data or (for obscure reasons having to do with computer hardware) as “double” or “double-precision” data.
- **Character data**, which represents text. Examples include names of months, or categories (such as the name of a city or country). This kind of data is also referred to as “string data”.
- **factor data**, which represents variables that can only take on certain discrete values. R treats factor data as a kind of augmented character data.

The difference between character data and factor data is that factor data has an explicit set of allowed values and has an integer number associated with each of those values.

For instance, if I have a factor variable with the allowed values “up” and “down”, then I could not assign it a value “left” or “right”, whereas a character variable can be assigned any arbitrary text, such as “second star to the right and straight on til morning.”

There are two kinds of factors: ordered and unordered. The difference is that ordered factors have a specific order, so you can say that one value comes before or after another (or is greater than or less than another), whereas unordered factors don’t have any natural ordering.

Examples of ordered variables might be the months of the year, or the days of the week, or a grouping like “small”, “medium”, “large”, or “bad”, “fair”, “good”.

Examples of unordered variables might be lists of states, gender, religion, cities, sports teams, or other descriptive characteristics that don’t have a natural order to them.

- **date and time data**, which represents calendar dates, times of day, or a combination, such as 9:37 PM on January 17, 1984.

For the most part, R handles different data types sensibly so you don’t need to worry about them, but sometimes when R is reading data in from files, or when you want to convert one kind of variable to another, you will need to think about these.

The most common cases where you will need to think about this is when you are reading data in from files. Sometimes it is ambiguous whether to treat something from a file as character data, numerical data, or a date. In such cases, you may need to give R guidance about how to interpret data. The functions for reading data in from files, such as `read_csv` and `read_table` allow you

to specify whether a given column of data in a table is integer, double-precision (floating point), character, date, etc.

R also provides functions for converting data. The `as.character` function takes data that might be character, factor, or numeric, and represent it as text (characters).

`as.numeric` or `as.integer` will allow you to convert a character variable to a number. For instance, `as.numeric("3.14")` converts a text variable “3.14” into a numeric variable 3.14.

`as.integer` is very useful when we want to convert an ordered factor to an integer that corresponds to the order of that value. For instance, if I have an ordered factor `f` with legal values corresponding to the months of the year (“Jan”, “Feb”, ..., “Dec”), if `f` has the value Mar, then `as.integer(f)` will have the integer value 3.

Vectors, Lists, Data Frames, and Tibbles.

In statistics, you generally don’t just work with one number at a time, but with collections of numbers. R provides many ways to work with collections of numbers.

Vectors

The simplest is a **vector**. A vector is a collection of values that are all of the same kind: a collection of integers, a collection of floating point values, a collection of character values, a collection of factor values, etc.

You specify vectors like this: `x = c(1, 2, 5, 9, 3, 4, 2, 7, 5)`. You can access elements of vectors by indexing their position within the vector, so `x[3]` will be 5 and `x[4]` will be 9.

You can also give the elements of a vector names: `ages = c(Sam = 27, Ben = 20, Sarah = 25, Deborah = 31)` allows you to use `ages["Ben"]`, which will be 20.

All of the elements of a vector have to be the same kind, so `x = c(1, 2, "three")` will not allow the vector to mix numbers and characters and R will transform all of the values to character. The result is "1", "2", "three", and `x[1] + x[2]` will give an error because R doesn’t know how to add two character variables. However, `as.numeric(x[1]) + as.numeric(x[2])` will yield 3.

Lists

Lists are a lot like vectors, but they can contain different kinds of variables. They can even contain lists and vectors. `x = list(1, 2, "three", list(4, 5, 6))` has four elements. The first two are the numbers 1 and 2; the third is the character string “three”, and the fourth is the list (4,5,6).

Just as we can have named vectors, we can have named lists: `ages = list(Sam = 27, Ben = 20, Sarah = 25, Deborah = 31)`. There is a nice shortcut to getting the elements of a named list, using a dollar sign: `ages$Sam` is 27.

Data frames and tibbles

We will not use lists very much in this class. We will use vectors a little bit, but what we will use *a lot* are tables of data. You have probably worked a lot with spreadsheets and other data analysis tools that organize data in a table with rows and columns. This is a very natural and common way to work.

R provides a structure called a `data.frame` for working with tabular data, but the package `tidyverse` introduces an improved version of the `data.frame` called a `tibble` (think of it as a kind of data table).

`data.frames` and `tibbles` have rows and columns. A row represents a set of quantities, such as measurements or observations, that go together in some way. Different rows in a `tibble` represent different sets of these quantities.

For instance, if I am measuring the height and weight of a number of people, then I would have a row for each person and each row would have a column for the person's name or identity code, a column for their height, and a column for their weight.

If I am measuring the average temperature and average precipitation for a number of cities, then I would have a column for the city, a column for the temperature, and a column for the precipitation.

Each column of a `data.frame` or `tibble` should correspond to a specific kind of data (integer, floating point, character, factor, date, etc.). A column is a kind of vector, so it has to obey the restrictions that apply to vectors.

To get some experience with `tibbles`, let's load a couple of data sets that I have prepared. If you have cloned the directory for this document (from https://github.com/gilligan-ees-3310/lab_02_documentation), you can load the datasets, which contain daily weather summaries for Nashville and Chicago, by running the code below:

```
nashville_weather = readRDS('_data/nashville_weather.Rds')
chicago_weather = readRDS('_data/chicago_weather.Rds')
```

Here is an example of the first few rows of a `tibble` with weather data for Nashville from 1948–2017:

```
head(nashville_weather)
```

```
## # A tibble: 6 x 6
##   id      date      prcp  tmin  tmax location
##   <chr>   <date>   <dbl> <dbl> <dbl> <chr>
## 1 USW00013897 1948-01-01  67.8  -1.1  18.3 Nashville, TN
## 2 USW00013897 1948-01-02   0    -1.7   1.7 Nashville, TN
## 3 USW00013897 1948-01-03   0    -3.3   9.4 Nashville, TN
## 4 USW00013897 1948-01-04   0     1.7   8.9 Nashville, TN
## 5 USW00013897 1948-01-05   0    -2.2   8.9 Nashville, TN
## 6 USW00013897 1948-01-06   0    -1.1  10.6 Nashville, TN
```

There are 6 columns: the weather station ID for the Nashville Airport, the date of the measurement, the daily precipitation (in millimeters), the daily minimum and maximum temperatures (Celsius), and the name of the location.

The tibble also shows the kind of variable that each column represents: `id` and `location` are character data, `date` is Date data, and `prcp`, `tmin`, and `tmax` are double-precision floating point data (i.e., real numbers).

In some ways, a tibble or `data.frame` is like a named list of vectors, where each vector is a column and its name is the name of the column. We can access individual columns using the dollar sign, just as with regular named lists:

```
precipitation = nashville_weather$prcp  
  
head(precipitation)
```

```
## [1] 67.8  0.0  0.0  0.0  0.0  0.0
```

RStudio has a nice feature that lets you examine a tibble or `data.frame` as though it were a spreadsheet. To examine one of R's built-in data sets, which has data on hurricanes in the Atlantic from 1975–2015: `View(dplyr::storms)`

Some other useful functions:

- You can get a list of the names of a named vector, a named list, or the columns of a tibble or `data.frame` with the `names` function: `names(x)`, where `x` is a vector, list, tibble, or `data.frame`.
- You can get the length of a vector or list with the `length` function, and you can get the number of rows and columns in a tibble using the `dim` function:

```
x = c(1, 2, 3, 4, 5)  
print("Length of x is")
```

```
## [1] "Length of x is"
```

```
print(length(x))
```

```
## [1] 5
```

```
print("Dimensions of dplyr::storms is ")
```

```
## [1] "Dimensions of dplyr::storms is "
```

```
dim(dplyr::storms)
```

```
## [1] 10010    13
```

That's 10010 rows and 13 columns. You can also get just the number of rows or the number of columns with `nrow()` and `ncol()`.

The Tidyverse

The “tidyverse” is a collection of packages written by Hadley Wickham to make it easy to work with data frames. Wickham developed an improved kind of data frame that has features that are lacking in the basic R `data.frame`, and he developed a collection of tools for manipulating, analyzing, and graphing data from tibbles and regular `data.frames`.

To use the tidyverse, we need to load the package using R's `library` function. If tidyverse is not installed on your computer, you will get an error message and you will have to run `install.packages("tidyverse")` before you can proceed.

When you load tidyverse, it automatically loads a bunch of useful packages for manipulating and analyzing data: `tibble`, `dplyr`, `tidyr`, `purrr`, `readr`, and `ggplot2`.

If you have the `pacman` package installed, it can help you avoid these error messages: after you load `pacman` with `library(pacman)`, then you can load other packages using `p_load(tidyverse)` (you can substitute any other package name for “tidyverse”): `pacman` will first see whether you have that package on your computer; if you do, `pacman` will load it, and if you don't `pacman` will install the package from the Comprehensive R Archive Network (CRAN) and then load it.

In the code below, I will also load the `lubridate` package, which is part of tidyverse but is not loaded automatically when you load tidyverse. `lubridate` provides useful functions for working with dates, which will come in handy as we work with the weather data.

```
library(tidyverse)
library(lubridate)
# alternately, I could do the following:
# library(pacman)
# p_load(tidyverse, lubridate)
```

One part of the tidyverse is the package `dplyr`, which has many useful tools for modifying and manipulating tibbles:

- `select` lets you choose a subset of columns from a tibble
- `rename` lets you rename columns
- `filter` lets you choose a subset of rows from a tibble
- `arrange` lets you sort the rows with respect to the values of different columns

- `mutate` lets you modify the values of columns or add new columns
- `summarize` lets you generate summaries of columns (e.g., the mean, maximum, or minimum value of that column)
- `group_by` and `ungroup` let you perform calculations with grouping (e.g., in combination with `summarize`, you can group by year to produce separate summaries for each year)
- `bind_rows` to combine multiple tibbles that have the same kinds of columns by stacking one above the other.

There is a lot more, but these functions will be enough to keep us busy for now and they will allow us to do some powerful analysis.

Let's start with `select`: You can select columns to keep or columns to delete.

Here are the first few rows of `nashville_weather`

```
head(nashville_weather)
```

```
## # A tibble: 6 x 6
##   id          date      prcp  tmin  tmax location
##   <chr>      <date>    <dbl> <dbl> <dbl> <chr>
## 1 USW00013897 1948-01-01  67.8  -1.1  18.3 Nashville, TN
## 2 USW00013897 1948-01-02   0    -1.7   1.7 Nashville, TN
## 3 USW00013897 1948-01-03   0    -3.3   9.4 Nashville, TN
## 4 USW00013897 1948-01-04   0     1.7   8.9 Nashville, TN
## 5 USW00013897 1948-01-05   0    -2.2   8.9 Nashville, TN
## 6 USW00013897 1948-01-06   0    -1.1  10.6 Nashville, TN
```

Let's get rid of the `id` column, since we don't really care about the ID number, that meteorological agencies use to identify the weather station. And let's get rid of the `location` column because we know that the data set is from Nashville, so having that information repeated on each row is a waste of space. To do this, we just call `select`, specifying the tibble or `data.frame` to operate on, and then give a list of columns to eliminate, with a minus sign in front of each:

```
x = select(nashville_weather, -id, -location)
head(x)
```

```
## # A tibble: 6 x 4
##   date      prcp  tmin  tmax
##   <date>    <dbl> <dbl> <dbl>
## 1 1948-01-01  67.8  -1.1  18.3
## 2 1948-01-02   0    -1.7   1.7
## 3 1948-01-03   0    -3.3   9.4
## 4 1948-01-04   0     1.7   8.9
## 5 1948-01-05   0    -2.2   8.9
## 6 1948-01-06   0    -1.1  10.6
```


Alternately, instead of telling select which columns to get rid of, we can tell it which columns to keep:

```
x = select(nashville_weather, date, prcp, tmin, tmax)
head(x)
```

```
## # A tibble: 6 x 4
##   date      prcp  tmin  tmax
##   <date>    <dbl> <dbl> <dbl>
## 1 1948-01-01  67.8  -1.1  18.3
## 2 1948-01-02   0    -1.7   1.7
## 3 1948-01-03   0    -3.3   9.4
## 4 1948-01-04   0     1.7   8.9
## 5 1948-01-05   0    -2.2   8.9
## 6 1948-01-06   0    -1.1  10.6
```

We can specify a range of consecutive columns by giving the first and last with a colon between them:

```
x = select(nashville_weather, date:tmax)
head(x)
```

```
## # A tibble: 6 x 4
##   date      prcp  tmin  tmax
##   <date>    <dbl> <dbl> <dbl>
## 1 1948-01-01  67.8  -1.1  18.3
## 2 1948-01-02   0    -1.7   1.7
## 3 1948-01-03   0    -3.3   9.4
## 4 1948-01-04   0     1.7   8.9
## 5 1948-01-05   0    -2.2   8.9
## 6 1948-01-06   0    -1.1  10.6
```

This is a general R thing: we can specify a range of numbers in a similar way:

```
1:10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

rename lets us rename columns:

```
x = rename(nashville_weather, weather_station = id, city = location)
head(x)
```

```
## # A tibble: 6 x 6
##   weather_station date      prcp  tmin  tmax city
##   <chr>          <date>    <dbl> <dbl> <dbl> <chr>
## 1 USW00013897  1948-01-01  67.8  -1.1  18.3 Nashville, TN
## 2 USW00013897  1948-01-02   0    -1.7   1.7 Nashville, TN
## 3 USW00013897  1948-01-03   0    -3.3   9.4 Nashville, TN
## 4 USW00013897  1948-01-04   0     1.7   8.9 Nashville, TN
## 5 USW00013897  1948-01-05   0    -2.2   8.9 Nashville, TN
## 6 USW00013897  1948-01-06   0    -1.1  10.6 Nashville, TN
```

filter lets us select only rows that match a condition:

```
x = filter(nashville_weather, year(date) > 2015 & tmax < 0)
head(x)
```

```
## # A tibble: 6 x 6
##   id      date      prcp  tmin  tmax location
##   <chr>   <date>    <dbl> <dbl> <dbl> <chr>
## 1 USW00013897 2016-01-18   0    -11   -3.8 Nashville, TN
## 2 USW00013897 2016-01-19   0    -9.9  -2.1 Nashville, TN
## 3 USW00013897 2016-01-23   0    -8.2  -1   Nashville, TN
## 4 USW00013897 2016-02-09   1.3   -6    -2.7 Nashville, TN
## 5 USW00013897 2016-02-10   0.5  -7.1  -0.5 Nashville, TN
## 6 USW00013897 2016-12-15   0    -6    -1   Nashville, TN
```

In the code above, I used the year function from the lubridate package to extract just the year from a date.

One thing that is important to know about making comparisons in filter expressions: to specify that two things are equal, you write == with two equal signs. A single equal sign is for assigning a value to a variable and two equal signs are for comparisons. You can also use “<=” for less than or equal to and “>=” for greater than or equal to.

```
x = filter(nashville_weather, date == ymd("2016-06-10"))
x
```

```
## # A tibble: 1 x 6
##   id      date      prcp  tmin  tmax location
##   <chr>   <date>    <dbl> <dbl> <dbl> <chr>
## 1 USW00013897 2016-06-10   0  15.6  33.3 Nashville, TN
```

In the code above, I used the ymd function from lubridate to translate a character value “2016-06-10” to the date value for June 10, 2016.

You can combine conditions in filter by using & to indicate “and” and | to indicate “or”.

We can sort the rows of a tibble or data.frame with the arrange function:

```
x = arrange(nashville_weather, desc(tmax), tmin)
head(x, 10)
```

```
## # A tibble: 10 x 6
##   id          date      prcp  tmin  tmax location
##   <chr>      <date>    <dbl> <dbl> <dbl> <chr>
## 1 USW00013897 2012-06-29      0    21.1  42.8 Nashville, TN
## 2 USW00013897 1952-07-27      0     20   41.7 Nashville, TN
## 3 USW00013897 1952-07-28      0    22.8  41.7 Nashville, TN
## 4 USW00013897 2012-06-30      0    26.7  41.7 Nashville, TN
## 5 USW00013897 1952-06-30   15.2    23.3  41.1 Nashville, TN
## 6 USW00013897 2007-08-16      0    24.4  41.1 Nashville, TN
## 7 USW00013897 2012-06-28      0    17.8  40.6 Nashville, TN
## 8 USW00013897 1988-07-08      0     20   40.6 Nashville, TN
## 9 USW00013897 1954-09-05      0    20.6  40.6 Nashville, TN
## 10 USW00013897 2012-07-06      0    22.2  40.6 Nashville, TN
```

This sorts the rows in descending order of tmax (i.e., so the largest values are at the top), and where multiple rows have the same value of tmax, then it sorts them in ascending order of tmin. Observe the three rows with tmax = 41.7, the two rows where tmax = 41.1, and the four rows where tmax = 40.6.

We can use mutate to modify the values of columns or to create new columns. The temperatures in Nashville weather are in Celsius and the precipitation is in millimeters. Let's convert these to Fahrenheit and inches, respectively, and then let's create a trange column that will have the difference between the maximum and minimum temperature:

```
x = mutate(nashville_weather, prcp = prcp / 25.4, tmin = tmin * 9./5. + 32,
           tmax = tmax * 9./5. + 32, trange = tmax - tmin)
head(x)
```

```
## # A tibble: 6 x 7
##   id          date      prcp  tmin  tmax location      trange
##   <chr>      <date>    <dbl> <dbl> <dbl> <chr>      <dbl>
## 1 USW00013897 1948-01-01   2.67  30.0  64.9 Nashville, TN  34.9
## 2 USW00013897 1948-01-02    0    28.9  35.1 Nashville, TN   6.12
## 3 USW00013897 1948-01-03    0    26.1  48.9 Nashville, TN  22.9
## 4 USW00013897 1948-01-04    0    35.1  48.0 Nashville, TN  13.0
## 5 USW00013897 1948-01-05    0    28.0  48.0 Nashville, TN  20.0
## 6 USW00013897 1948-01-06    0    30.0  51.1 Nashville, TN  21.1
```

Summaries are useful for finding averages and extreme values. Let's find the maximum and minimum temperatures and the most extreme rainfall in the whole data set:

```
x = summarize(nashville_weather, prcp.max = max(prcp),
              tmin.min = min(tmin), tmax.max = max(tmax))
x
```

```
## # A tibble: 1 x 3
##   prcp.max tmin.min tmax.max
##   <dbl>    <dbl>    <dbl>
## 1      NA      NA      NA
```

nashville_weather has 25446 rows, but summarize reduces it to a single summary row.

You can use summary to generate multiple summary quantities from a column:

```
x = summarize(nashville_weather, prcp.max = max(prcp), prcp.min = min(prcp))
x
```

```
## # A tibble: 1 x 2
##   prcp.max prcp.min
##   <dbl>    <dbl>
## 1      NA      NA
```

We can also generate grouped summaries:

```
x = ungroup(summarize(group_by(nashville_weather, year(date)),
                          prcp.max = max(prcp), prcp.tot = sum(prcp)))
head(x)
```

```
## # A tibble: 6 x 3
##   `year(date)` prcp.max prcp.tot
##   <dbl>    <dbl>    <dbl>
## 1     1948     67.8    1179.
## 2     1949     78.2    1324.
## 3     1950     84.6    1633
## 4     1951     97.3    1484.
## 5     1952    118.     1011.
## 6     1953     47.8    1050.
```

This provides the maximum one-day precipitation and the total annual precipitation for each year

Note how difficult it is to read that grouped summary expression: the group_by function is inside summarize, which is inside ungroup.

The tidyverse offers us a much nicer way to put these kind of complicated expressions together using what it calls the “pipe” operator, %>%. The pipe operator chains operations together, taking the result of the one on the left and inserting it into the one on the right.

We can use the pipe operator to rewrite the expression above as

```
x = nashville_weather %>% group_by(year(date)) %>%
  summarize(prcp.max = max(prcp), prcp.tot = sum(prcp)) %>%
  ungroup()
head(x)
```

```
## # A tibble: 6 x 3
##   `year(date)` prcp.max prcp.tot
##         <dbl>    <dbl>    <dbl>
## 1      1948      67.8    1179.
## 2      1949      78.2    1324.
## 3      1950      84.6    1633
## 4      1951      97.3    1484.
## 5      1952     118.     1011.
## 6      1953      47.8     1050.
```

Now the expression is easier to read: First we group `nashville_weather` by year, then we summarize it by calculating the maximum daily precipitation and the yearly total for each year, and finally, after we summarize we ungroup.

You can combine any set of the tidyverse functions using the pipe operator, so you could select columns, filter rows, mutate the values of columns, and summarize, using the `%>%` pipe operator to connect all of the different operations in sequence.

The final dplyr command we're going to look at is `bind_rows`, which lets us combine tibbles:

```
weather = bind_rows(nashville_weather, chicago_weather)
```

This creates a single tibble that has all of the rows from `nashville_weather` on top and all the rows from `chicago_weather` on the bottom. Because the two tibbles have the same columns, the columns are matched up.

We can operate on this combined tibble:

```
weather_summary = weather %>%
  mutate(year = year(date), t.range = tmax - tmin) %>%
  group_by(year, location) %>%
  summarize(prcp.max = max(prcp), prcp.tot = sum(prcp),
            t.range.max = max(t.range)) %>%
  ungroup() %>%
  arrange(year, location)
tail(weather_summary)
```

```
## # A tibble: 6 x 5
##   year location      prcp.max prcp.tot t.range.max
```

```
##      <dbl> <chr>                <dbl>    <dbl>        <dbl>
## 1  2015 Chicago, IL             98.3    1169.         18.3
## 2  2015 Nashville, TN          59.9    1291.         23.9
## 3  2016 Chicago, IL            50.8    1067.         23.9
## 4  2016 Nashville, TN          52.8    1086.         22.2
## 5  2017 Chicago, IL            NA        NA            NA
## 6  2017 Nashville, TN          NA        NA            NA
```

Re-shaping tibbles: gather and spread

Sometimes you want to gather many columns in a data table, data frame, or tibble into a single column. For instance, consider this data frame, which was read in from a spreadsheet of global temperatures produced at NASA:


```
giss_zonal = readRDS('_data/giss_zonal.Rds')
head(giss_zonal)
```

```
## # A tibble: 6 x 9
##   year x64n_90n x44n_64n x24n_44n equ_24n x24s_equ x44s_24s x64s_44s x90s_64s
##   <int>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1  1880   -1.05   -0.51   -0.26   -0.21   -0.14   -0.05    0.05    0.65
## 2  1881   -0.93   -0.5    -0.15    0.03    0.06   -0.07   -0.07   0.570
## 3  1882   -1.51  -0.290  -0.09   -0.09   -0.07    0.02    0.04    0.6
## 4  1883   -0.4    -0.62  -0.22   -0.24   -0.19   -0.02    0.07   0.48
## 5  1884   -1.36  -0.62  -0.41   -0.23   -0.2    -0.19   -0.02   0.63
## 6  1885   -1.19  -0.68  -0.42   -0.09   -0.23   -0.31   -0.15   0.8
```

The tibble presents the average temperatures for different bands of latitude: 64°N–90°N, 44°N–64°N, 24°N–44°N, Equator–24°N, and the same for the Southern Hemisphere.

If we wanted to plot all of these, we could do something like this:

```
ggplot(giss_zonal, aes(x = year)) +
  geom_line(aes(y = x64n_90n, color = "64N-90N")) +
  geom_line(aes(y = x44n_64n, color = "44N-64N")) +
  geom_line(aes(y = x24n_44n, color = "24N-44N")) +
  geom_line(aes(y = equ_24n, color = "EQU-24N")) +
  geom_line(aes(y = x24s_equ, color = "24S-EQU")) +
  geom_line(aes(y = x44s_24s, color = "44S-24S")) +
  geom_line(aes(y = x64s_44s, color = "64S-44S")) +
  geom_line(aes(y = x90s_64s, color = "90S-64S")) +
  labs(x = "Year", y = "Temperature anomaly")
```



This is a big mess. It would be hard to clean up the appearance, and would require a lot of retyping if we decided to group the data into different bands of latitude.

We can do this much more easily with the gather function:

```
bands = names(giss_zonal) # column names of the tibble
bands = bands[-1] # drop the first column ("year")
labels = c("64N-90N", "44N-64N", "24N-44N", "EQU-24N",
           "24S-EQU", "44S-24S", "64S-44S", "90S-64S")
tidy_zonal = giss_zonal %>%
  gather(key = latitude, value = anomaly, -year)

head(tidy_zonal)
```

```
## # A tibble: 6 x 3
##   year latitude anomaly
##   <int> <chr>      <dbl>
## 1  1880 x64n_90n   -1.05
## 2  1881 x64n_90n   -0.93
## 3  1882 x64n_90n   -1.51
## 4  1883 x64n_90n   -0.4
## 5  1884 x64n_90n   -1.36
```

```
## 6 1885 x64n_90n -1.19
```


Now we can clean up the latitude column a bit to make it more friendly for human readers:

```
tidy_zonal = tidy_zonal %>%  
  mutate(latitude = ordered(latitude, levels = bands,  
                             labels = labels)) %>%  
  # ^^^ the previous line converts the latitude band into an ordered factor  
  # where the order is the order of the original columns. This will prevent  
  # R from sorting them alphabetically when it makes the legend for the plot.  
  # The "labels" parameter then changes the names from the somewhat cryptic  
  # original column name to something a human can read easily.  
  arrange(year, latitude)  
  
head(tidy_zonal)
```

```
## # A tibble: 6 x 3  
##   year latitude anomaly  
##   <int> <ord>      <dbl>  
## 1  1880 64N-90N    -1.05  
## 2  1880 44N-64N    -0.51  
## 3  1880 24N-44N    -0.26  
## 4  1880 EQU-24N    -0.21  
## 5  1880 24S-EQU    -0.14  
## 6  1880 44S-24S    -0.05
```

Now let's plot it:

```
ggplot(tidy_zonal, aes(x = year, y = anomaly, color = latitude)) +  
  geom_line() +  
  labs(x = "Year", y = "Temperature anomaly")
```

The code for making the plot was a lot simpler, and by using an ordered factor, we could control the order of the latitude bands in the legend, which now appear in a sensible order. It is much easier to look at this graph and quickly recognize that the far northern latitudes (64N-90N, and to a lesser extent 44N-64N) are warming up much faster than the rest of the planet.

Back in 1967, one of the first global climate models predicted that global warming due to greenhouse gases would cause the far northern latitudes to warm up much faster than the rest of the planet. This data confirms that prediction.

We also see that the Southern Hemisphere has warmed much less than the Northern. Think about why that might be.

We can also do the inverse of `gather` and spread one column of data and one “key” column into many columns, whose names are taken from the “key” column.

Let’s go back to the `weather_summary` tibble we made above:

```
head(weather_summary)
```

```
## # A tibble: 6 x 5
##   year location    prcp.max prcp.tot t.range.max
##   <dbl> <chr>         <dbl>    <dbl>    <dbl>
## 1  1928 Chicago, IL      NA      NA      NA
## 2  1929 Chicago, IL      NA      NA      NA
```

```
## 3 1930 Chicago, IL      NA      NA      NA
## 4 1931 Chicago, IL      NA      NA      NA
## 5 1932 Chicago, IL      NA      NA      NA
## 6 1933 Chicago, IL      NA      NA      NA
```

Let's set it up to make it easy to compare the annual precipitation of Nashville and Chicago:

```
x = weather_summary %>% select(year, location, prcp.tot) %>%
  spread(key = location, value = prcp.tot)
tail(x)
```

```
## # A tibble: 6 x 3
##   year `Chicago, IL` `Nashville, TN`
##   <dbl>      <dbl>      <dbl>
## 1 2012         801.        1165.
## 2 2013        1076.        1394.
## 3 2014        1322.        1286.
## 4 2015        1169.        1291.
## 5 2016        1067.        1086.
## 6 2017         NA          NA
```

Graphing Data

Here, we will look at the `ggplot2` package for plotting data. This package is automatically loaded when you load the `tidyverse` collection with `library(tidyverse)`. It follows a theory of making useful graphs of data called, “The Grammar of Graphics” (that’s where the “gg” comes from).

The idea is that a graph has several distinct parts, which come together:


- One or more *layers* of graphics. A layer consists of the following:
 - A *data table* with one or more columns, each corresponding to a different variable,
 - A *mapping* of different variables (columns) in the data table to different, *aesthetics* of the plot. Aesthetics are things like
 - * the *x* coordinate,
 - * the *y* coordinate,
 - * the *color* of the point or line,
 - * the *fill* color that is used to fill in areas, like the interior of a rectangle or circle.
 - * the *shape* of points (e.g., circle, square, triangle, cross, diamond, ...)
 - * the *size* of points and lines
 - * the *linetype* (e.g., solid, dashed, dotted, ...)
 - * and so forth ...
 - A *geometry* (point, line, box, etc.) that is used to draw the data

- A coordinate system (axes and legends)

There are some more aspects to the grammar of graphics, but we don't need them for what we're going to do.


A simple graph has just one layer:

```
ggplot(data = tidy_zonal, # the data
       # the mapping of variables to aesthetics
       mapping = aes(x = year, y = anomaly, color = latitude)
     ) +
  geom_point() + # the geometry
  labs(x = "Year", y = "Temperature Anomaly") # labels for the coordinates
```



We can also make a plot with the same data, but two layers:

```
ggplot(data = tidy_zonal, # the data
       # the mapping of variables to aesthetics
       mapping = aes(x = year, y = anomaly, color = latitude)
     ) +
  geom_point() + # the geometry of the first layer
  geom_line() + # the geometry of the second layer
  labs(x = "Year", y = "Temperature Anomaly") # labels for the coordinates
```



We can also use different mappings for different layers. Here, I add the shape aesthetic to the points to distinguish daily maximum and minimum temperatures.

Here, we run into an important issue with R: there are some days where the measurements of daily maximum or minimum temperature are missing for one of the weather stations. These missing values are represented in R with the special value NA, which means “not available”.

If you try to summarize a bunch of numbers (taking the maximum, minimum, average, etc. over the bunch) where some of them are missing, R’s default behavior is to report NA for the summary because you don’t know how the missing values should contribute to the summary (e.g., if you’re looking for the hottest or coldest day of the year, and don’t have measurements for several days, you don’t know whether those days were extremely hot or cold).

You can change this default behavior by using the `na.rm` (remove missing values) option of many summary functions, such as `max`, `min`, `mean` and `sd`. If we set `na.rm = TRUE`, as in the code below, then R will ignore the missing values and calculate the summaries based on the measurements we do have.

```
annual_extremes = weather %>% mutate(year = year(date)) %>%
  group_by(location, year) %>%
  # the na.rm = T below means to ignore missing values
  # if we don't put that in, then if any year has
  # a missing value for even one day, the tmax
```

```


# or tmin for that year will be recorded as
# NA (missing)
summarize(tmin = min(tmin, na.rm = T),
           tmax = max(tmax, na.rm = T)) %>%
ungroup()

ggplot(data = annual_extremes # the data
) +
  geom_point(aes(x = year, y = tmin, color = location, shape = "min")) +
  geom_point(aes(x = year, y = tmax, color = location, shape = "max")) +
  xlim(1990,2000) + # set the x-axis to go from 1990-2000
  labs(x = "Year", y = "Temperature Anomaly") # labels for the coordinates

```

```
## Warning: Removed 138 rows containing missing values (geom_point).
```

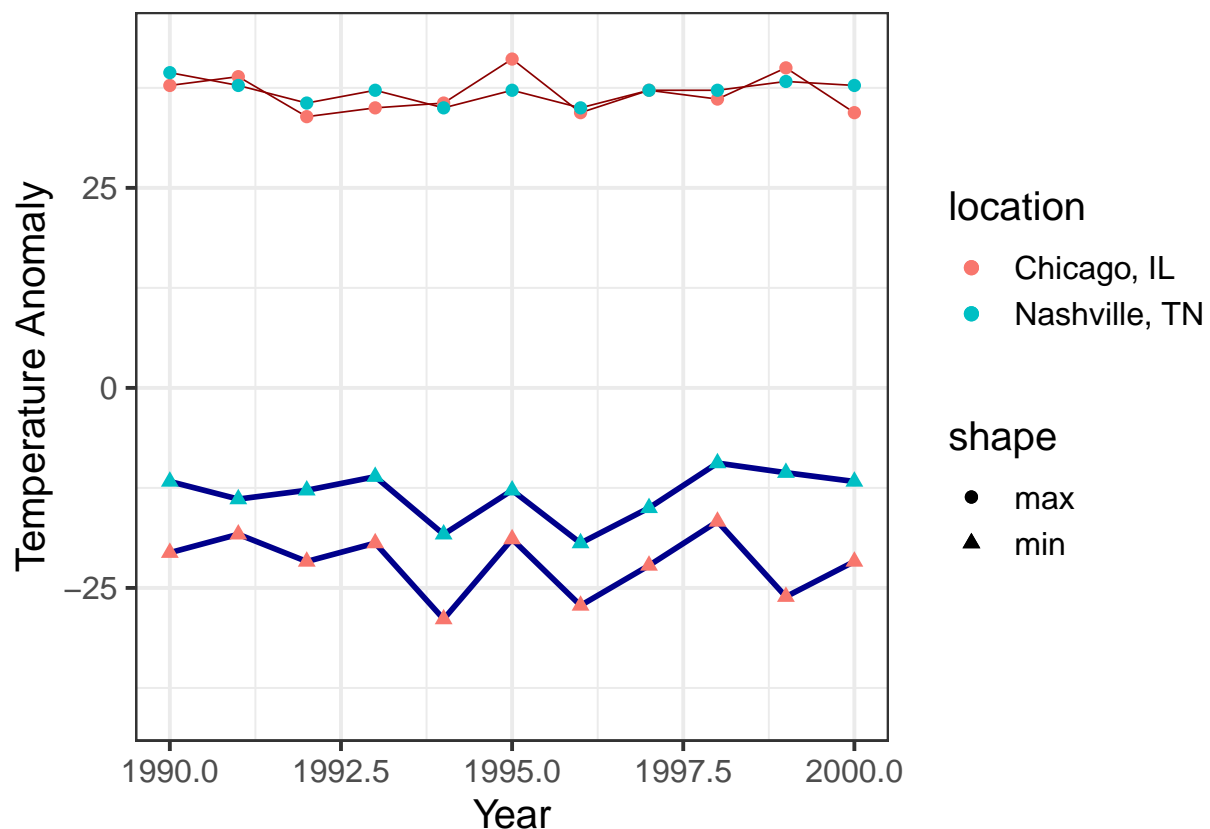
```
## Warning: Removed 138 rows containing missing values (geom_point).
```



Note that ggplot issued several harmless warnings to tell us that setting the limits of the x -axis the way we did. We can tell RMarkdown not to include those warnings in the document by adding “warning=FALSE” to the options for the chunk

If we want to specify aesthetics as having fixed values, we can specify them outside of the mapping. Here I specify the size and color of lines:

```
ggplot(data = annual_extremes # the data
) +
  geom_line(aes(x = year, y = tmin, group = location),
            color = "darkblue", size = 1) +
  geom_line(aes(x = year, y = tmax, group = location),
            color = "darkred", size = 0.3) +
  geom_point(aes(x = year, y = tmin, color = location, shape = "min"),
             size = 2) +
  geom_point(aes(x = year, y = tmax, color = location, shape = "max"),
             size = 2) +
  xlim(1990,2000) +
  labs(x = "Year", y = "Temperature Anomaly") # labels for the coordinates
```



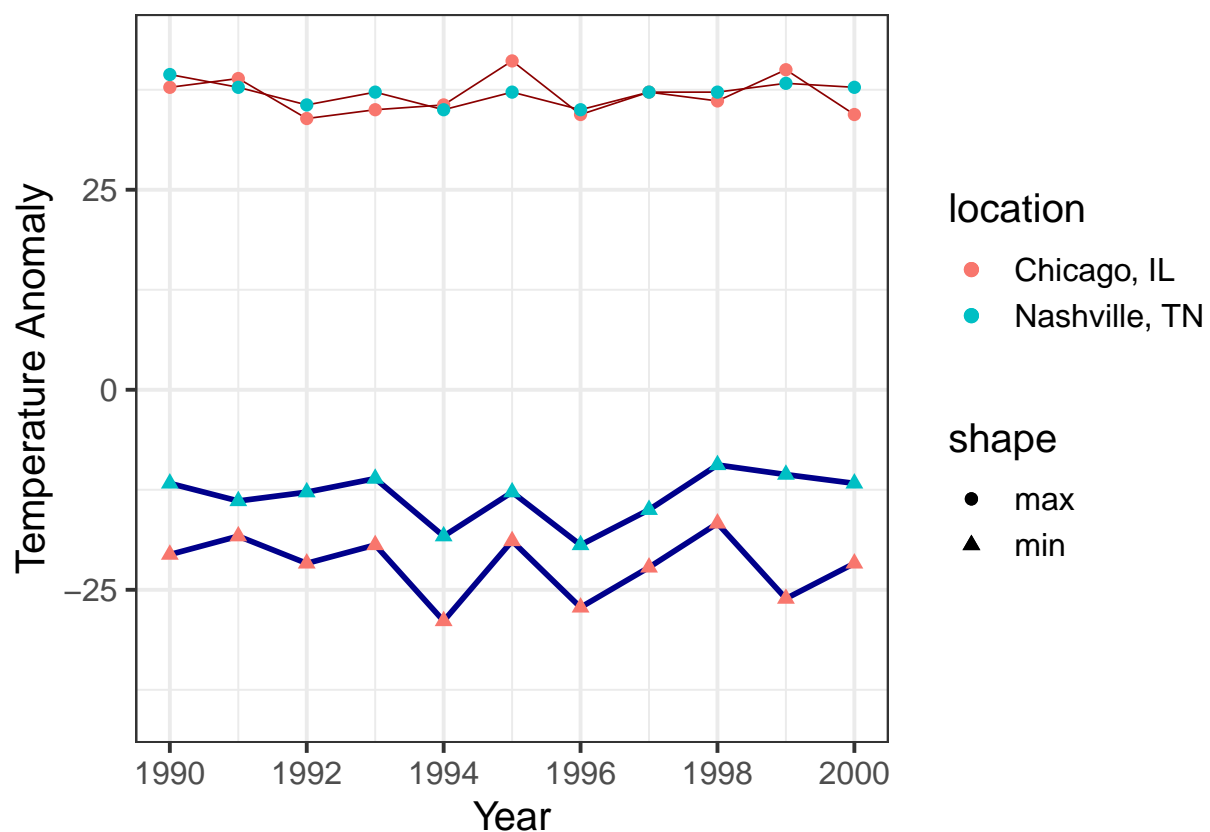
We can also take finer control of the axis formatting:

```
ggplot(data = annual_extremes # the data
) +
  geom_line(aes(x = year, y = tmin, group = location),
```

```

    color = "darkblue", size = 1) +
  geom_line(aes(x = year, y = tmax, group = location),
    color = "darkred", size = 0.3) +
  geom_point(aes(x = year, y = tmin, color = location, shape = "min"),
    size = 2) +
  geom_point(aes(x = year, y = tmax, color = location, shape = "max"),
    size = 2) +
  scale_x_continuous(limits=c(1990,2000),
    breaks = c(1990, 1992, 1994, 1996, 1998, 2000)) +
  # ^^^ the "breaks" parameter for an axis tells R where to put the labels
  labs(x = "Year", y = "Temperature Anomaly") # labels for the coordinates

```



And, of course, we could use `gather` to simplify this graph. `Gather` re-organizes the tibble with the data by combining the `tmin` and `tmax` columns into one column called `value`, and adds a new column that I call `Temperature`, which gets the names for the columns. This way, each temperature measurements gets its own row in the `value` columns, and each row has either `"tmax"` or `"tmin"` in the `Temperature` column to tell us what kind of temperature measurement it is.

An advantage of organizing the data in this way is that we can then tell `ggplot` to pick the aesthetics (like shape and size) for points and lines automatically, depending on the `Temperature` column for that row.

In the figure below, I also tell ggplot to use a special color palette instead of R's default colors. Here, I use the ColorBrewer palette, designed by Cynthia Brewer to be easier for people with different kinds of color blindness to see.

R has a number of different `scale_color_...` commands for choosing colors and many users have published custom color and styling packages that you can install and use to customize the way your plots look.

```
annual_extremes_gathered = annual_extremes %>%  
  gather(key = Temperature, value = value, -year, -location)  
  
ggplot(data = annual_extremes_gathered # the data  
  ) +  
  geom_line(aes(x = year, y = value, color = location, size = Temperature)) +  
  geom_point(aes(x = year, y = value, color = location, shape = Temperature),  
    size = 2) +  
  # set coordinates for "size"  
  scale_size_manual(values = c(tmax = 1.0, tmin = 0.1)) +  
  # Set coordinates for the x-axis.  
  scale_x_continuous(limits=c(1990,2000),  
    breaks = c(1990, 1992, 1994, 1996, 1998, 2000)) +  
  scale_color_brewer(palette = "Set1", name = "City") +  
  labs(x = "Year", y = "Temperature Anomaly") # labels for the coordinates
```